

Introduction to Deep Generative Models

Part 1

Deep Autoregressive Models

Fatemeh Saleh



Australian
National
University



Generative Models

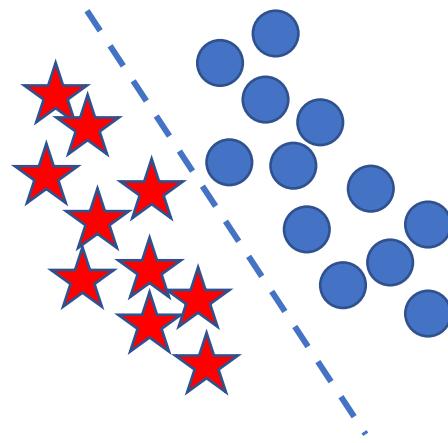
- Informally:
 - A model that can generate new data after learning from the dataset.
- More formally:
 - A generative model models the joint distribution $P(X, Y)$ of the observation X and the target Y .
 - A discriminative model models the conditional distribution $P(Y|X)$.

Discriminative versus Generative

- Discriminative Model
 - Tries to learn the discriminative information from the data
 - Example: Classify C1 vs C2 vs C3
 - Finds a good decision boundary by directly modeling conditional distribution $P(Y|X)$
 - Learns mappings from inputs to classes
- Generative Model
 - Tries to learn the distribution of the data
 - Models the distribution of inputs characteristic of the class
 - For classification, builds a model of $P(X|Y)$ and then applies Bayes Rule


Discriminative versus Generative

Discriminative Model



Learn $P(Y|X)$ directly
Logistic regression use a sigmoid function to estimate this directly

Generative Model




Learn $P(X|Y)$ and $P(Y)$

Learning What do “stars” look like?

The class Prior
 $p(y=0)$ $p(y=1)$

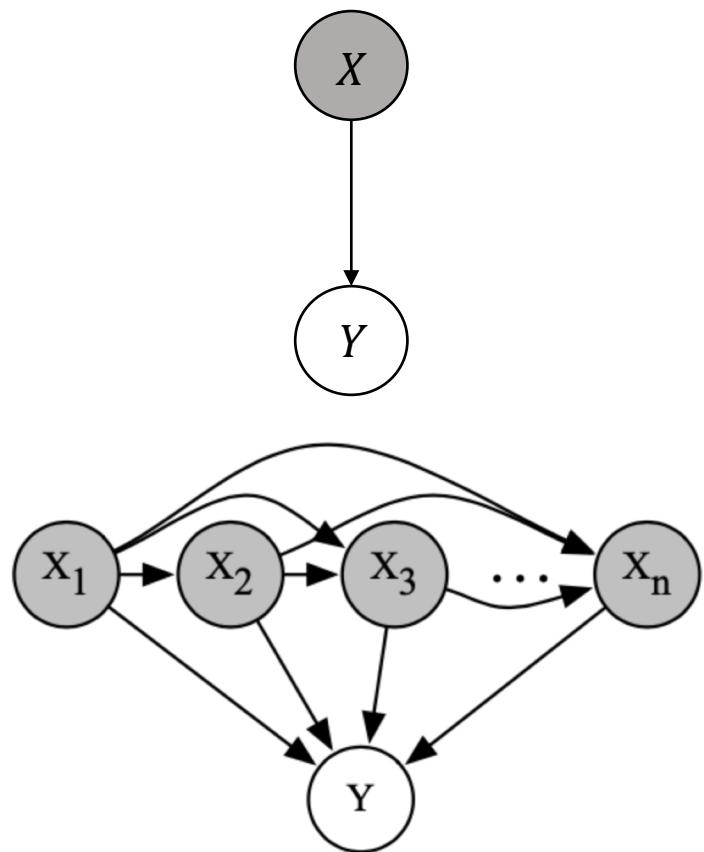
Discriminative versus Generative

- Suppose that: $P(X|Y)$ and $P(Y)$
- New example x
- We can then compute $P(y = 1|x) = \frac{p(x|y = 1)p(y=1)}{p(x)}$ (Bayes Rule)
- where $p(x) = \sum_y p(x,y) = p(x|y = 1)p(y = 1) + p(x|y = 0)p(y = 0)$

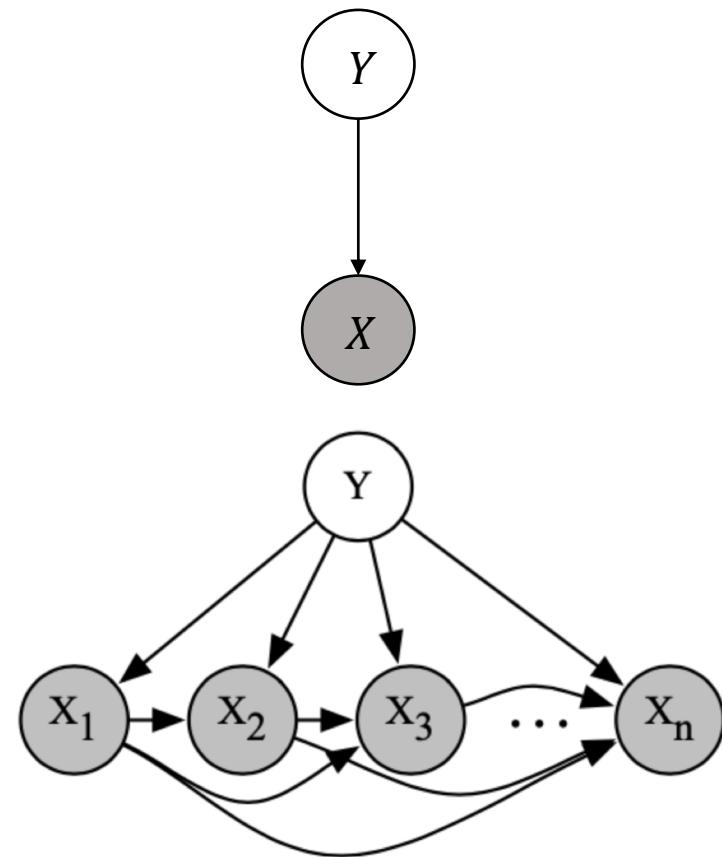


Discriminative versus Generative

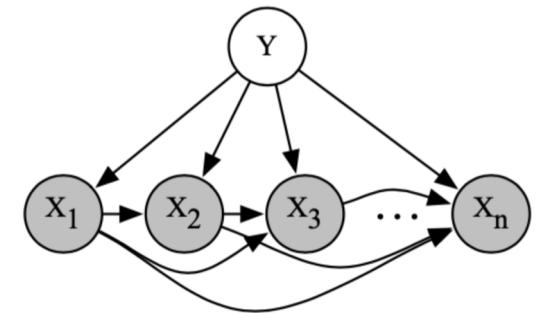
Discriminative Model



Generative Model

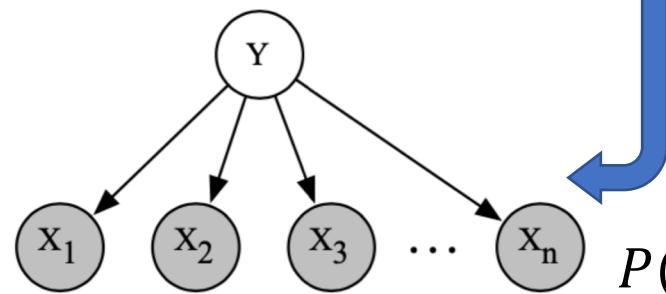


Generative Model (naive Bayes)



$$P(y, x_1, x_2, \dots, x_n) = P(y)P(x_1|y)P(x_2|y, x_1) \dots P(x_n|y, x_1, \dots, x_{n-1})$$

$$x_1 \perp x_2 \perp \dots \perp x_n | y$$

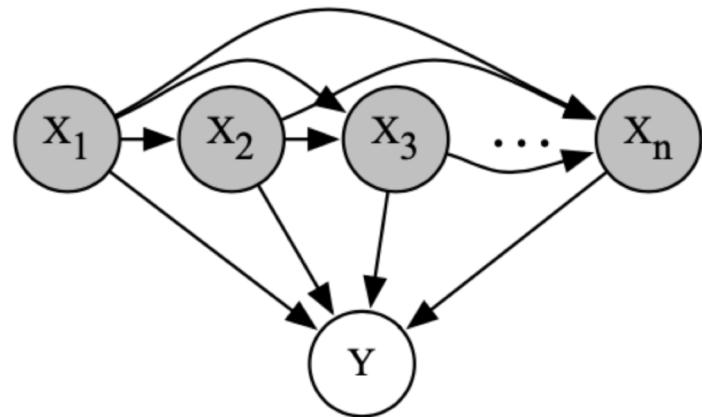


$$P(y, x_1, x_2, \dots, x_n) = P(y)P(x_1|y)P(x_2|y) \dots P(x_n|y) = P(y) \prod_{i=1}^n P(x_i|y)$$

Estimate parameters from training data. Predict with Bayes rule:

$$P(Y = 1|x_1, x_2, \dots, x_n) = \frac{P(Y = 1) \prod_{i=1}^n P(x_i|Y = 1)}{\sum_{y=\{0,1\}} P(Y = y) \prod_{i=1}^n P(x_i|Y = y)}$$

Discriminative Model (logistic regression)



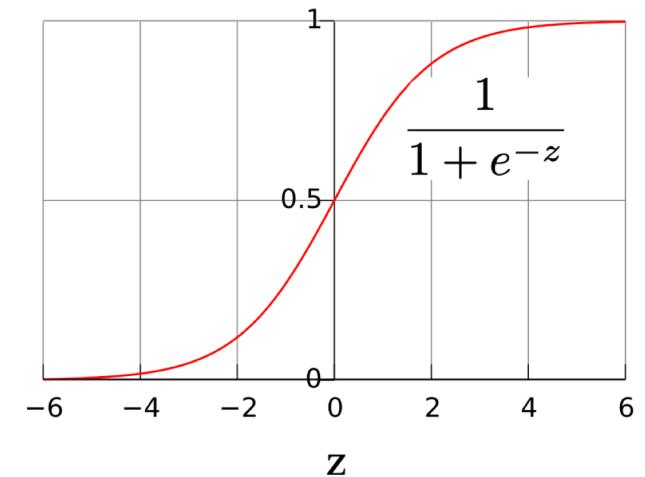
$$P(Y = 1|X; \alpha) = f(X, \alpha)$$

$$z(\alpha, X) = \alpha_0 + \sum_{i=1}^n \alpha_i x_i$$

$$P(Y = 1|X; \alpha) = \sigma(z(X, \alpha))$$

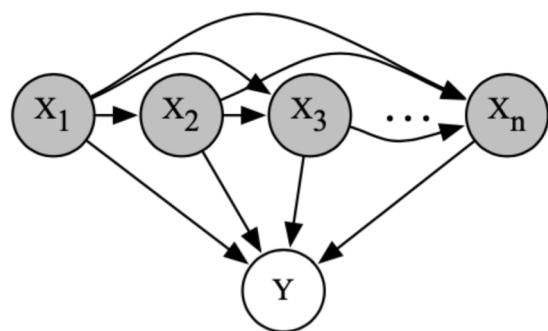
where

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

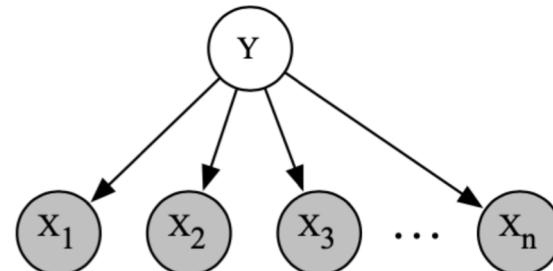


Discriminative versus Generative

Discriminative (logistic regression)



Generative (Naïve Bayes)



Discriminative models are powerful

- Logistic model does not assume $x_i \perp x_{i-1} | y$, unlike naïve Bayes
- This can make a big difference in many applications
 - For example, in spam classification, let $x_1 = 1["bank" \text{ in email}]$ and $x_2 = 1["account" \text{ in email}]$
 - Regardless of whether spam, these always appear together, i.e. $x_1 = x_2$
 - Learning in naïve Bayes results in $p(x_1|y) = p(x_2|y)$, thus naïve Bayes double counts the evidence
 - Learning with logistic regression sets $\alpha_1 = 0$ or $\alpha_2 = 0$, in effect ignoring it

Generative models are still very useful

- Using a conditional model is only possible when X is always observed

Generative Models

- Given a training set of examples, e.g., images of cats:



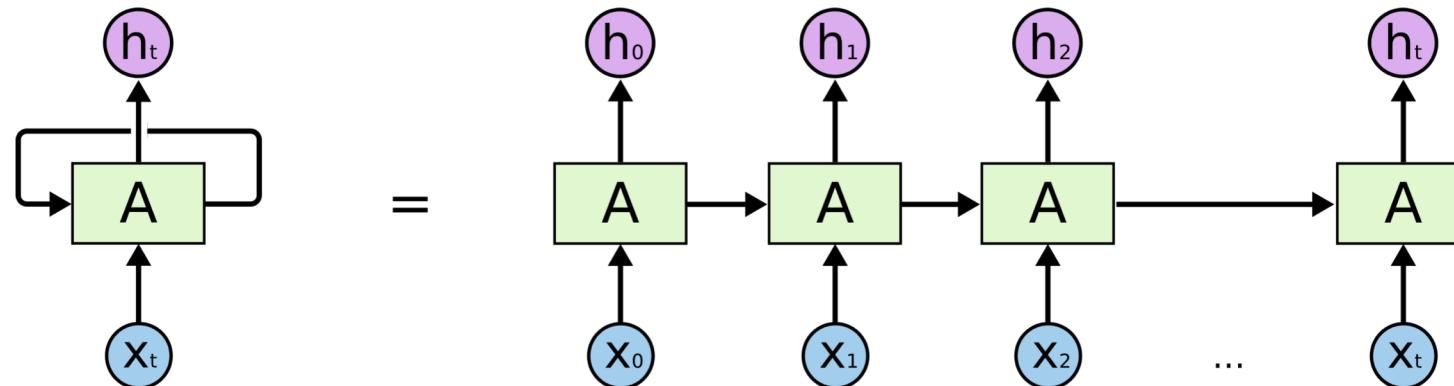
- We want to learn a probability distribution $p(x)$ over images x for
 - Generation:** If we sample $x_{new} \sim p(x)$, x_{new} should look like a cat (sampling)
 - Density estimation:** $p(x)$ should be high if x looks like a cat, and low otherwise
 - Unsupervised representation learning:** The model should be able to learn what these images have in common, e.g., ears, tail, etc. (features)

Autoregressive Models

Autoregressive Models

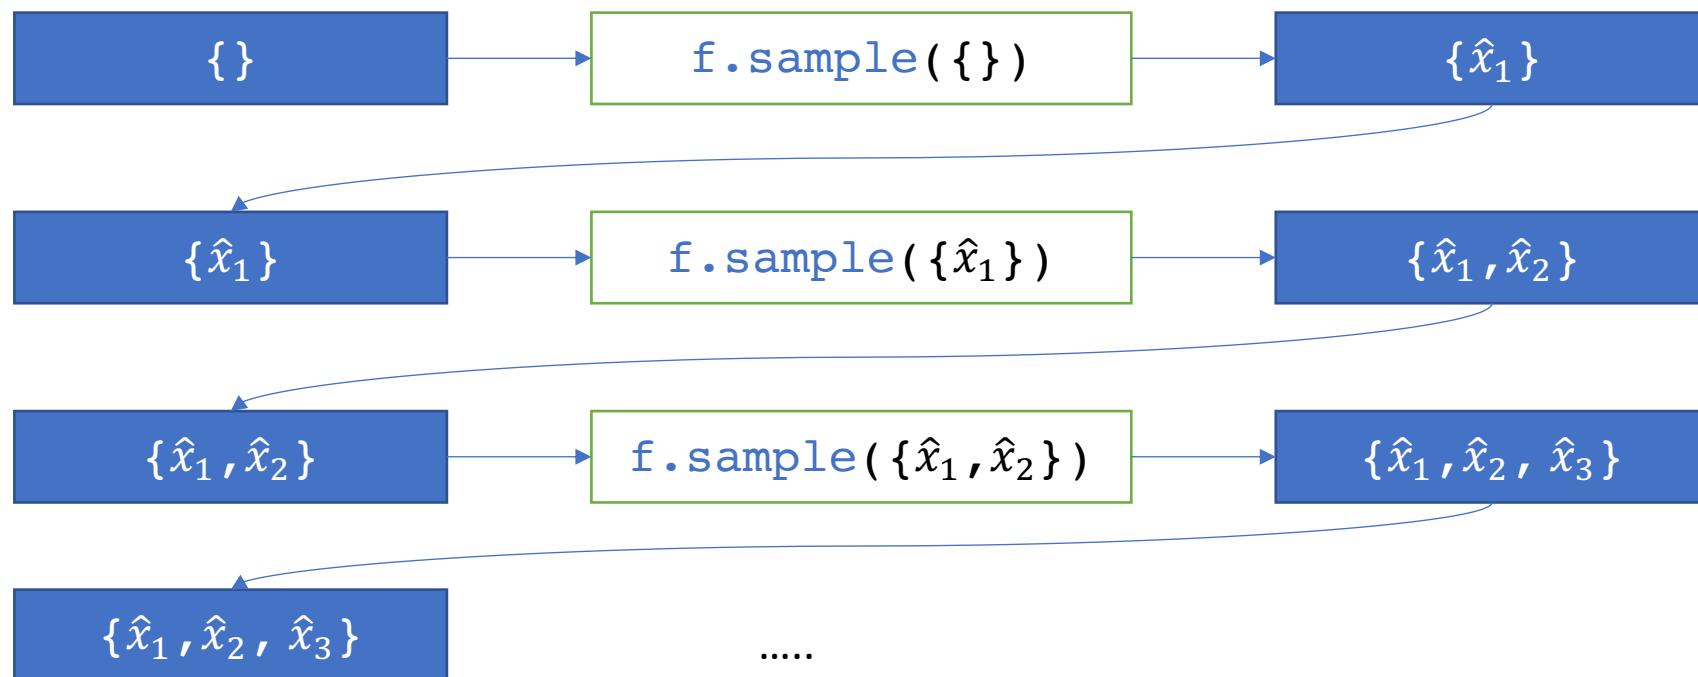
- Definition:
 - Given the observation as sequences x_1, x_2, \dots, x_T , we decompose the likelihood into a product of conditional distributions:

$$p(x_1, x_2, \dots, x_T) = \prod_{t=1}^T p(x_t | x_1, \dots, x_{t-1})$$



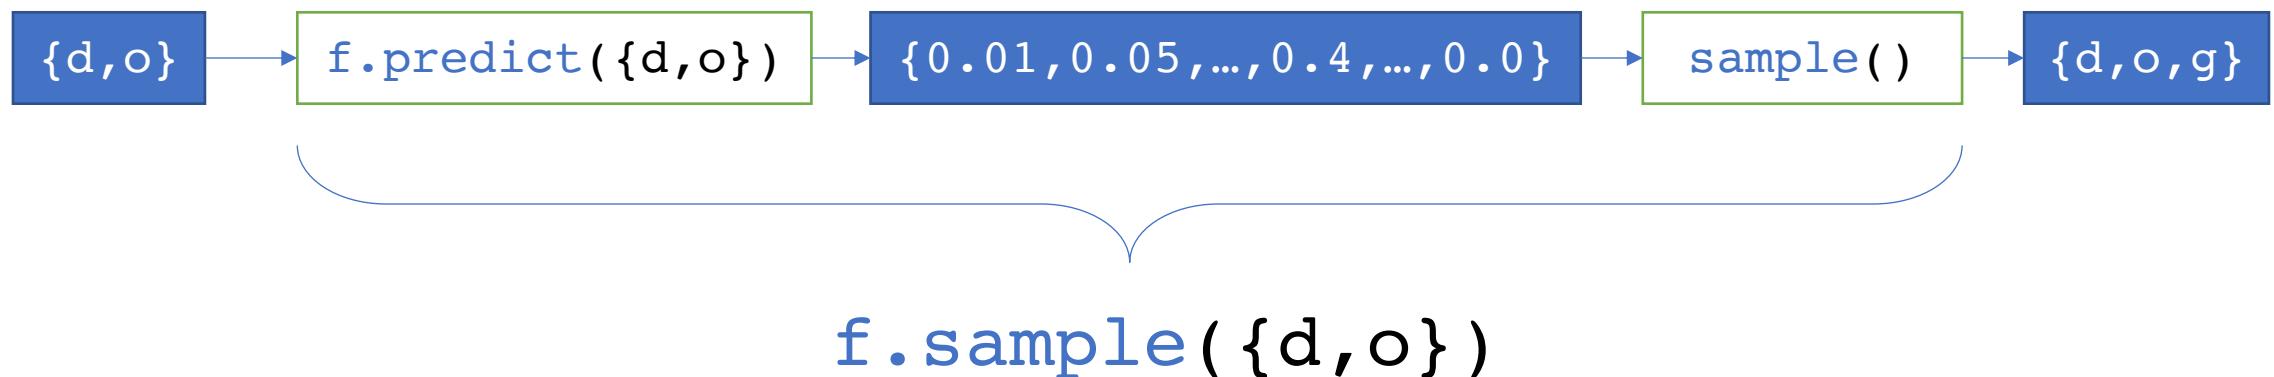
Autoregressive Models

- Given such a model `f` and a sampling function `sample`, the generative process for a full sequence is:

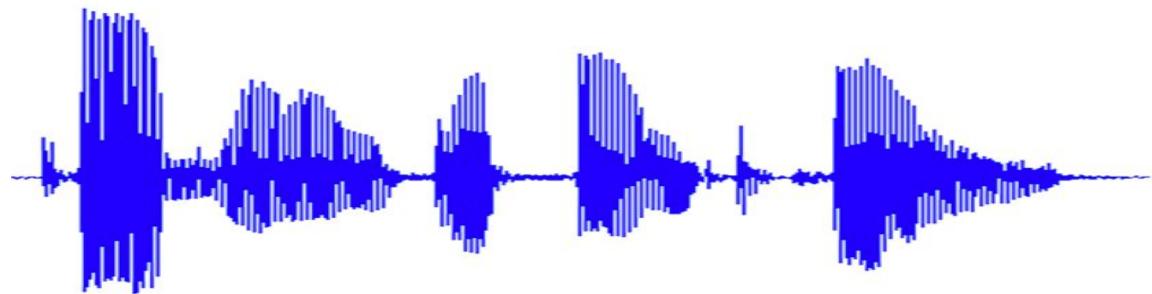
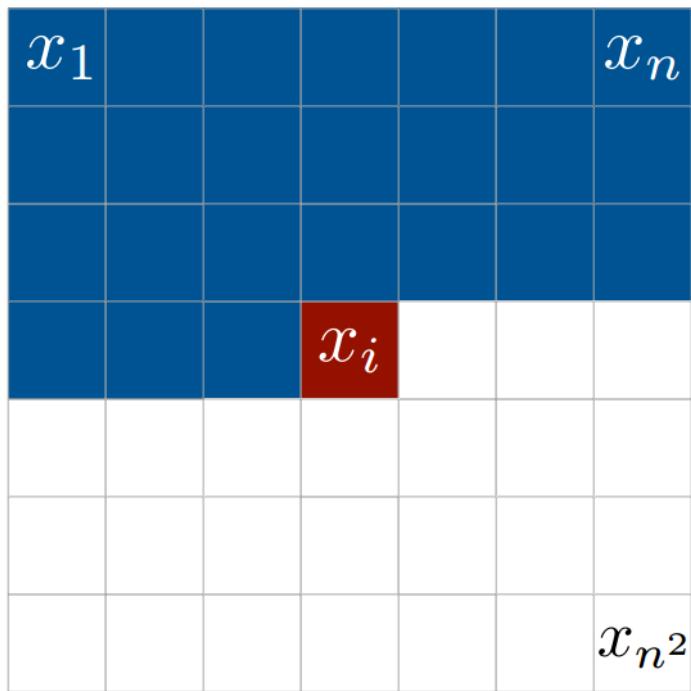


Autoregressive Models

- How `f.sample(.)` works?
 - Consider the example of character generation, where we only have C possible options to choose from.
 - The model maps the d dimensional data to C dimensional probability distribution and then samples from the estimated distribution.




Autoregressive Models




Autoregressive Models

- One solution is to use *causal convolutions*.
 - How to create causal convolutions?




Autoregressive Models

- One solution is to use *causal convolutions*.
 - How to create causal convolutions?



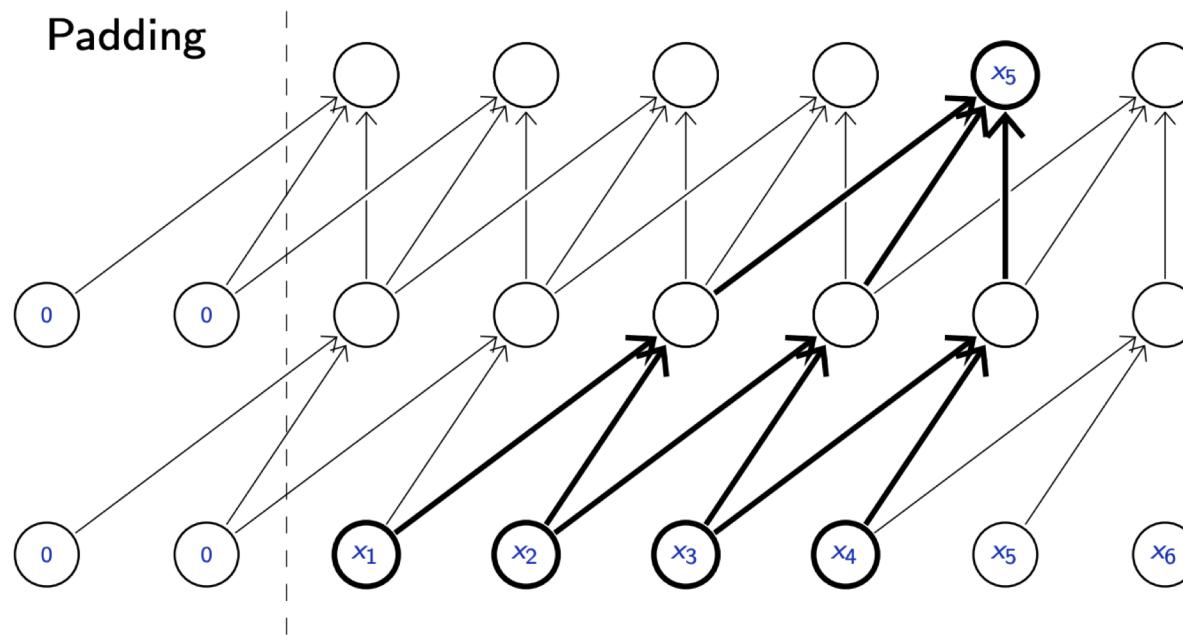
Autoregressive Models

- One solution is to use *causal convolutions*.
 - How to create causal convolutions?




Autoregressive Models

- One solution is to use *causal convolutions*.
 - How to create causal convolutions?




Autoregressive Models

- One solution is to use *causal convolutions*.
 - How to create causal convolutions?


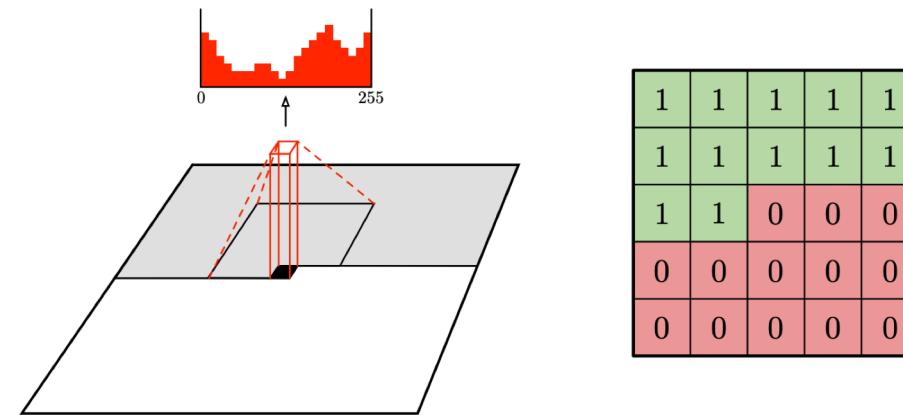


Autoregressive Models



Autoregressive Models

- Example: PixelCNN



An Example of Autoregressive Generation

Human Motion Prediction

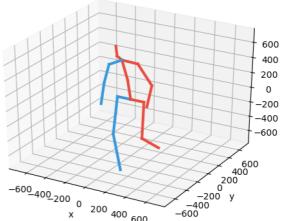
Autoregressive Models

- Let's look at an implementation of autoregressive human motion prediction. This code contains
 - Definition of the model
 - training data
 - **training** script
 - **sampling** script
- *Note: Here, we regress the future human pose as opposed to estimating a probability distribution and then sampling a pose from that.*

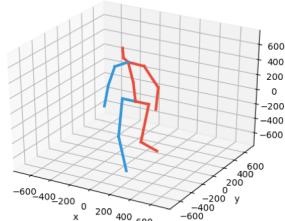
Motion Prediction (data)

- A sequence of pose: data in the form of $B \times T \times J$
 - Input during training: $\text{data}[:, :-1, :]$

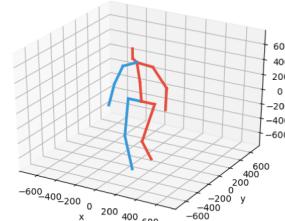
Walking frame:1 [Observation]



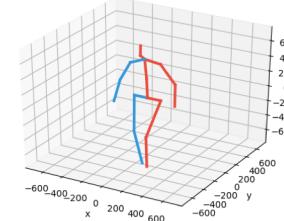
Walking frame:3 [Observation]



Walking frame:5 [Observation]

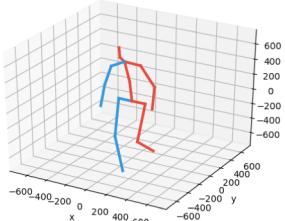


Walking frame:7 [Observation]

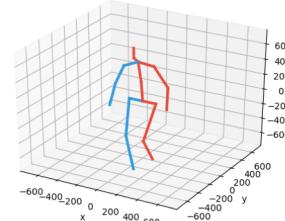


- Output during training: $\text{data}[:, 1:, :]$

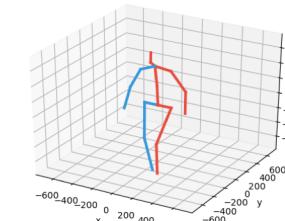
Walking frame:3 [Observation]



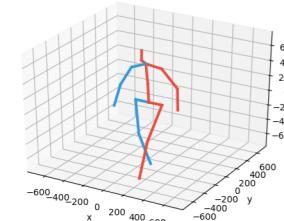
Walking frame:5 [Observation]



Walking frame:7 [Observation]




Walking frame:9 [Observation]



Motion Prediction (model)

```
1 import torch
2 from torch import nn
3
4
5 class MotionModel(nn.Module):
6     def __init__(self, data_dim=96, embedding_dim=256, hidden_dim=256):
7         super(MotionModel, self).__init__()
8
9         self.encoder = nn.Sequential(
10             nn.Linear(data_dim, embedding_dim),
11             nn.ReLU())
12         self.rnn = nn.GRU(embedding_dim, hidden_dim, num_layers=2)
13         self.decoder = nn.Sequential(
14             nn.Linear(hidden_dim, embedding_dim),
15             nn.ReLU(),
16             nn.Linear(embedding_dim, data_dim))
```



Motion Prediction (training and forward)

```
7 def train(model, trainloader, test_seq):
8     criterion = nn.MSELoss()
9     optimizer = optim.Adam(model.parameters(), lr=0.001)
10
11    for epoch in range(100):
12        model.train()
13        for batch_idx, (_, _, data) in enumerate(trainloader):
14            data = data.cuda()
15            inputs = data[:, :-1, :]
16            outputs = data[:, 1:, :]
17
18            optimizer.zero_grad()
19            predictions = model(inputs) // Line 19
20            loss = criterion(predictions,
21                            outputs)
22            loss.backward()
23            nn.utils.clip_grad_norm_(model.parameters(), 0.1)
24            optimizer.step()
25
```

```
18     def single_forward(self, x, h=None):
19
20         x_e = self.encoder(x)
21         x, h = self.rnn(x_e, h)
22         x = x + x_e.clone()
23         return self.decoder(x), h
24
25     def forward(self, data):
26
27         rec, h = self.single_forward(data)
28
29         return rec
```

Motion Prediction (generation)

```
30     def sample(self, observation, T=50):
31         output = []
32         obs_len = observation.shape[1]
33
34         rec, h = self.single_forward(observation[:, 0:1, :])
35
36         for t in range(1, obs_len):
37             rec, h = self.single_forward(observation[:, t:t+1, :], h)
38
39         output.append(rec)
40
41         for t in range(T-1):
42             rec, h = self.single_forward(rec, h)
43             output.append(rec)
44
45         return output
```

Forwarding the observation

Motion Prediction (generation)

```
30     def sample(self, observation, T=50):
31         output = []
32         obs_len = observation.shape[1]
33
34         rec, h = self.single_forward(observation[:, 0:1, :])
35
36         for t in range(1, obs_len):
37             rec, h = self.single_forward(observation[:, t:t+1, :], h)
38
39         output.append(rec)
40
41         for t in range(T-1):
42             rec, h = self.single_forward(rec, h)
43             output.append(rec)
44
45         return output
```

We use the groundtruth motion as the input to the model

Motion Prediction (generation)

```
30     def sample(self, observation, T=50):
31         output = []
32         obs_len = observation.shape[1]
33
34         rec, h = self.single_forward(observation[:, 0:1, :])
35
36         for t in range(1, obs_len):
37             rec, h = self.single_forward(observation[:, t:t+1, :], h)
38
39         output.append(rec)
40
41         for t in range(T-1):
42             rec, h = self.single_forward(rec, h)
43             output.append(rec)
44
45         return output
```

Given the last observation, we generate the first future motion

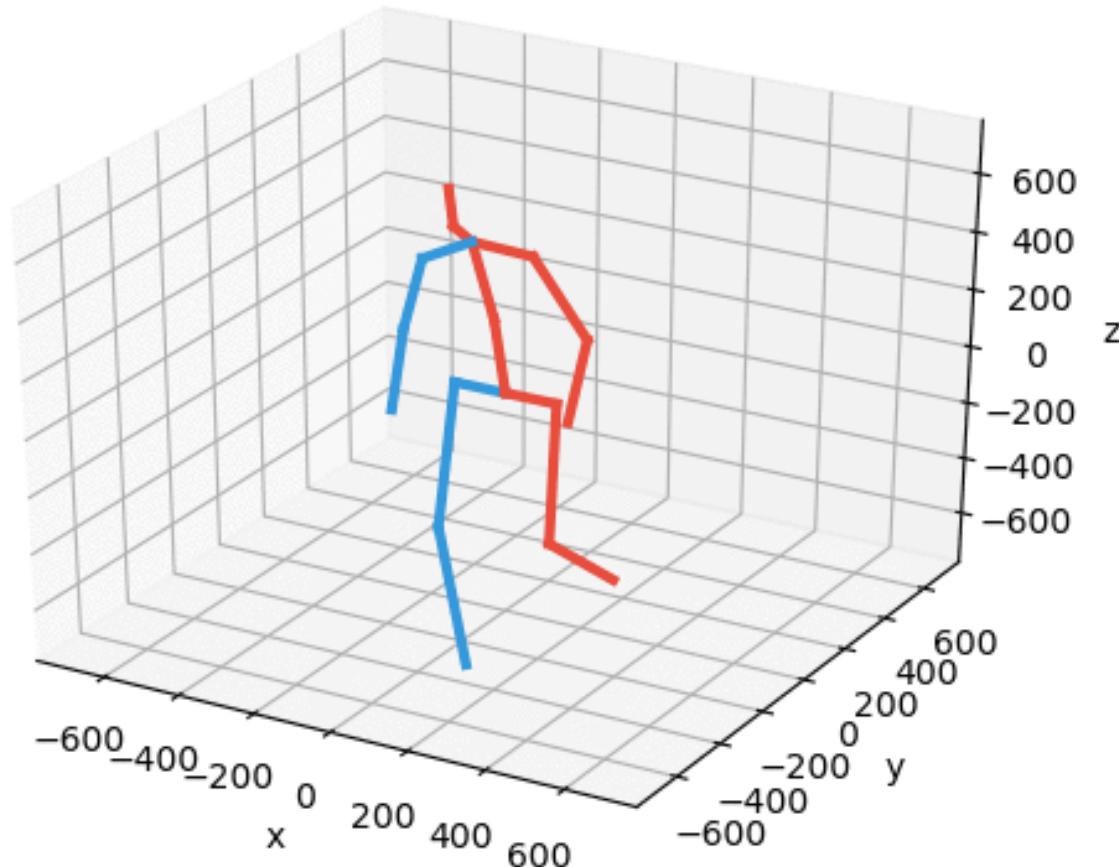
Motion Prediction (generation)

```
30     def sample(self, observation, T=50):
31         output = []
32         obs_len = observation.shape[1]
33
34         rec, h = self.single_forward(observation[:, 0:1, :])
35
36         for t in range(1, obs_len):
37             rec, h = self.single_forward(observation[:, t:t+1, :], h)
38
39         output.append(rec)
40
41         for t in range(T-1):
42             rec, h = self.single_forward(rec, h)
43             output.append(rec)
44
45     return output
```

Autoregressively, we generate future poses; given last prediction, we generate the next motion

Motion Prediction (results)

Walking frame:1 [Observation]



Summary

- An overview of generative versus discriminative models
- Autoregressive models
 - Pros:
 - Easy to sample from
 - Easy to compute the likelihood
 - Natural to train via maximum likelihood
 - Cons:
 - Slow to generate new data
 - Does not directly learn unsupervised representations of the data.
- Materials:
 - <https://deepgenerativemodels.github.io/>
 - <https://fleuret.org/ee559/>
 - Van Den Oord, et.al. "Pixel recurrent neural networks." ICML. 2016.